

Mawlana Bhashani Science and Technology University



Lab-Report

Lab Report No : 04

Lab Report Name : SDN Controllers and Mininet

Course Title : Computer Networks Lab

Submitted by

Name: Moniruzzaman & Mst.Zakia Sultana

ID: IT-18007 & IT-18027

3rd year 2nd semester

Session: 2017-2018

Dept. of ICT

MBSTU.

Submitted To

Nazrul Islam

Assistant Professor

Dept. of ICT

MBSTU.

Theory:

What is iPerf?

- iPerf is a tool for active measurements of the maximum achievable bandwidth on IP networks. It supports tuning of various parameters related to timing, buffers and protocols (TCP, UDP, SCTP with IPv4 and IPv6). For each test it reports the bandwidth, loss, and other parameters.

Software-Defined vs. Traditional Networking:

The key difference between traditional and software-defined networking is how SDNs handle data packets. In a traditional network, the way a switch handles an incoming data packet is written into its firmware. Most switches — particularly those used in commercial data centers rather than enterprise environments — respond to and route all packets the same way. SDN provides admins with granular control over the way switches handle data, giving them the ability to automatically prioritize or block certain types of packets. This, in turn, allows for greater efficiency without the need to invest in expensive, application-specific network switches.

Benefits of Software-Defined Networking: There are several benefits to the more advanced level of control afforded by implementing SND in a multi-tenant network environment:

- **Automation:** SND allows for automation of complex operational tasks that make networks faster, more efficient and easier to manage.
- **Increased uptime:** SDN has proven effective in reducing deployment and configuration errors that can lead to service disruptions.
- **Less drain on resources:** SDN gives administrators control over how their routers and switches operate from a single, virtual workflow. This frees up key staff to focus on more important tasks.
- **Better visibility:** With SDN, system administrator's gain improved visibility into overall network function, allowing them to allocate resources more effectively.
- **Cost savings:** SND can lead to significant overall costs savings. It also reduces the amount of spending required on infrastructure by allowing data centers to get the most use of their existing devices.

Controller:

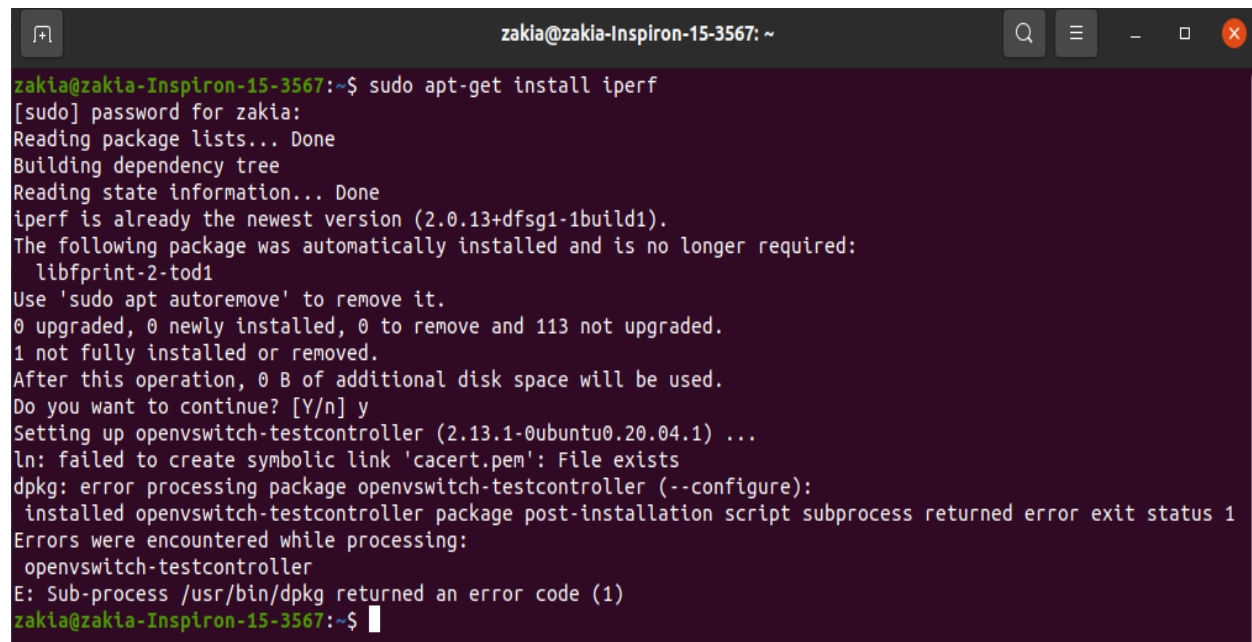
OVS-testcontroller is a simple OpenFlow controller that manages any number of switches over the OpenFlow protocol, causing them to function as L2 MAC-learning switches or hubs. It is suitable for initial testing of OpenFlow networks.

Ryu is a component-based software defined networking framework. Ryu provides software components with well-defined API that make it easy for developers to create new network management and control applications. Ryu supports various protocols for managing network devices, such as OpenFlow, Netconf, OF-config, etc. About OpenFlow, Ryu supports fully 1.0, 1.2, 1.3, 1.4, 1.5 and Nicira Extensions. All of the code is freely available under the Apache 2.0 license

Mininet:

Mininet creates a realistic virtual network, running real kernel, switch and application code, on a single machine (VM, cloud or native) Because you can easily interact with your network using the Mininet CLI (and API), customize it, share it with others, or deploy it on real hardware, Mininet is useful for development, teaching, and research. Mininet is also a great way to develop, share, and experiment with OpenFlow and Software-Defined Networking systems.

Install iperf:



```
zakia@zakia-Inspiron-15-3567: ~  
zakia@zakia-Inspiron-15-3567:~$ sudo apt-get install iperf  
[sudo] password for zakia:  
Reading package lists... Done  
Building dependency tree  
Reading state information... Done  
iperf is already the newest version (2.0.13+dfsg1-1build1).  
The following package was automatically installed and is no longer required:  
  libfprint-2-tod1  
Use 'sudo apt autoremove' to remove it.  
0 upgraded, 0 newly installed, 0 to remove and 113 not upgraded.  
1 not fully installed or removed.  
After this operation, 0 B of additional disk space will be used.  
Do you want to continue? [Y/n] y  
Setting up openvswitch-testcontroller (2.13.1-0ubuntu0.20.04.1) ...  
ln: failed to create symbolic link 'cacert.pem': File exists  
dpkg: error processing package openvswitch-testcontroller (--configure):  
  installed openvswitch-testcontroller package post-installation script subprocess returned error exit status 1  
Errors were encountered while processing:  
  openvswitch-testcontroller  
E: Sub-process /usr/bin/dpkg returned an error code (1)  
zakia@zakia-Inspiron-15-3567:~$
```

Install Mininet:

```
zakia@zakia-Inspiron-15-3567: ~  
zakia@zakia-Inspiron-15-3567:~$ sudo apt-get install mininet  
Reading package lists... Done  
Building dependency tree  
Reading state information... Done  
mininet is already the newest version (2.2.2-Subuntu1).  
The following package was automatically installed and is no longer required:  
  libfprint-2-tod1  
Use 'sudo apt autoremove' to remove it.  
0 upgraded, 0 newly installed, 0 to remove and 113 not upgraded.  
1 not fully installed or removed.  
After this operation, 0 B of additional disk space will be used.  
Do you want to continue? [Y/n] y  
Setting up openvswitch-testcontroller (2.13.1-0ubuntu0.20.04.1) ...  
ln: failed to create symbolic link 'cacert.pem': File exists  
dpkg: error processing package openvswitch-testcontroller (--configure):  
  installed openvswitch-testcontroller package post-installation script subprocess returned error exit status 1  
Errors were encountered while processing:  
  openvswitch-testcontroller  
E: Sub-process /usr/bin/dpkg returned an error code (1)  
zakia@zakia-Inspiron-15-3567:~$
```

4. Exercises

4.1.1: Open a Linux terminal, and execute the command line `iperf --help`.

Provide four configuration options of `iperf`.

```
zakia@zakia-Inspiron-15-3567: ~  
zakia@zakia-Inspiron-15-3567:~$ iperf --help  
Usage: iperf [-s|-c host] [options]  
       iperf [-h|--help] [-v|--version]  
  
Client/Server:  
  -b, --bandwidth #[kmgKMG | pps]  bandwidth to send at in bits/sec or packets per second  
  -e, --enhancedreports             use enhanced reporting giving more tcp/udp and traffic information  
  -f, --format [kmgKMG]            format to report: Kbits, Mbits, KBytes, MBytes  
  -i, --interval #                  seconds between periodic bandwidth reports  
  -l, --len #[kmKM]                length of buffer in bytes to read or write (Defaults: TCP=128K, v4 UDP=1470, v6 U  
DP=1450)  
  -m, --print_mss                   print TCP maximum segment size (MTU - TCP/IP header)  
  -o, --output <filename>          output the report or error message to this specified file  
  -p, --port #                      server port to listen on/connect to  
  -u, --udp                         use UDP rather than TCP  
      --udp-counters-64bit          use 64 bit sequence numbers with UDP  
  -w, --window #[KM]               TCP window size (socket buffer size)  
  -z, --realtime                    request realtime scheduler  
  -B, --bind <host>[:<port>][%<dev>] bind to <host>, ip addr (including multicast address) and optional port a  
nd device  
  -C, --compatibility               for use with older versions does not sent extra msgs  
  -M, --mss #                      set TCP maximum segment size (MTU - 40 bytes)  
  -N, --nodelay                     set TCP no delay, disabling Nagle's Algorithm  
  -S, --tos #                       set the socket's IP_TOS (byte) field  
  
Server specific:  
  -s, --server                      run in server mode  
  -t, --time #                      time in seconds to listen for new connections as well as to receive traffic (default  
t not set)  
      --udp-histogram #,#          enable UDP latency histogram(s) with bin width and count, e.g. 1,1000=1(ms),1000(bi  
ns)  
  -B, --bind <ip>[%<dev>]          bind to multicast address and optional device  
  -H, --ssm-host <ip>              set the SSM source, use with -B for (S,G)  
  -U, --single_udp                  run in single threaded UDP mode
```

```

-D, --daemon          run the server as a daemon
-V, --ipv6_domain     Enable IPV6 reception by setting the domain and socket to AF_INET6 (Can receive on
both IPV4 and IPV6)

Client specific:
-c, --client <host>   run in client mode, connecting to <host>
-d, --dualtest         Do a bidirectional test simultaneously
--ipg                 set the the interpacket gap (milliseconds) for packets within an isochronous frame
--isochronous <frames-per-second>:<mean>,<stddev> send traffic in bursts (frames - emulate video traffi
c)
-n, --num #[kmgKMG]   number of bytes to transmit (instead of -t)
-r, --tradeoff         Do a bidirectional test individually
-t, --time #          time in seconds to transmit for (default 10 secs)
-B, --bind [<ip> | <ip:port>] bind ip (and optional port) from which to source traffic
-F, --fileinput <name> input the data to be transmitted from a file
-I, --stdin            input the data to be transmitted from stdin
-L, --listenport #     port to receive bidirectional tests back on
-P, --parallel #       number of parallel client threads to run
-R, --reverse          reverse the test (client receives, server sends)
-T, --ttl #           time-to-live, for multicast (default 1)
-V, --ipv6_domain     Set the domain to IPV6 (send packets over IPV6)
-X, --peer-detect      perform server version detection and version exchange
-Z, --linux-congestion <algo> set TCP congestion control algorithm (Linux only)

Miscellaneous:
-x, --reportexclude [CDMSV] exclude C(connection) D(data) M(multicast) S(settings) V(server) reports
-y, --reportstyle C    report as a Comma-Separated Values
-h, --help             print this message and quit
-v, --version          print version information and quit

[kmgKMG] Indicates options that support a k,m,g,K,M or G suffix
Lowercase format characters are 10^3 based and uppercase are 2^n based
(e.g. 1k = 1000, 1K = 1024, 1m = 1,000,000 and 1M = 1,048,576)

```

The TCP window size option can be set by the environment variable TCP_WINDOW_SIZE. Most other options can be set by an environment variable IPERF_<long option name>, such as IPERF_BANDWIDTH.

Source at <<http://sourceforge.net/projects/iperf2/>>

Report bugs to <iperf-users@lists.sourceforge.net>

zakia@zakia-Inspiron-15-3567:~\$

Exercise 4.1.2: Open two Linux terminals, and configure terminal-1 as client (iperf -c IPv4_server_address) and terminal-2 as server (iperf -s).

For terminal -1:

```

zakia@zakia-Inspiron-15-3567: ~
zakia@zakia-Inspiron-15-3567:~$ iperf -s
-----
Server listening on TCP port 5001
TCP window size: 128 KByte (default)
-----

```

For terminal -2:

```
zakia@zakia-Inspiron-15-3567: ~  
zakia@zakia-Inspiron-15-3567:~$ iperf -c 127.0.0.1 -u  
-----  
Client connecting to 127.0.0.1, UDP port 5001  
Sending 1470 byte datagrams, IPG target: 11215.21 us (kalman adjust)  
UDP buffer size: 208 KByte (default)  
-----  
[ 3] local 127.0.0.1 port 58362 connected with 127.0.0.1 port 5001  
read failed: Connection refused  
[ 3] WARNING: did not receive ack of last datagram after 1 tries.  
[ ID] Interval      Transfer    Bandwidth  
[ 3] 0.0-10.0 sec  1.25 MBytes 1.05 Mbits/sec  
[ 3] Sent 891 datagrams  
zakia@zakia-Inspiron-15-3567:~$
```

Exercise 4.1.3: Open two Linux terminals, and configure terminal-1 as client and terminal-2 as server for exchanging UDP traffic, which are the command lines? Which are the statistics are provided at the end of transmission?

```
zakia@zakia-Inspiron-15-3567: ~  
zakia@zakia-Inspiron-15-3567:~$ iperf -c 127.0.0.1 -u  
-----  
Client connecting to 127.0.0.1, UDP port 5001  
Sending 1470 byte datagrams, IPG target: 11215.21 us (kalman adjust)  
UDP buffer size: 208 KByte (default)  
-----  
[ 3] local 127.0.0.1 port 42980 connected with 127.0.0.1 port 5001  
read failed: Connection refused  
[ 3] WARNING: did not receive ack of last datagram after 1 tries.  
[ ID] Interval      Transfer    Bandwidth  
[ 3] 0.0-10.0 sec  1.25 MBytes 1.05 Mbits/sec  
[ 3] Sent 891 datagrams  
zakia@zakia-Inspiron-15-3567:~$
```

```
zakia@zakia-Inspiron-15-3567: ~  
zakia@zakia-Inspiron-15-3567:~$ iperf -s -u  
-----  
Server listening on UDP port 5001  
Receiving 1470 byte datagrams  
UDP buffer size: 208 KByte (default)  
-----
```

Exercise 4.1.4: Open two Linux terminals, and configure terminal-1 as client and terminal-2 as server for exchanging UDP traffic, with:

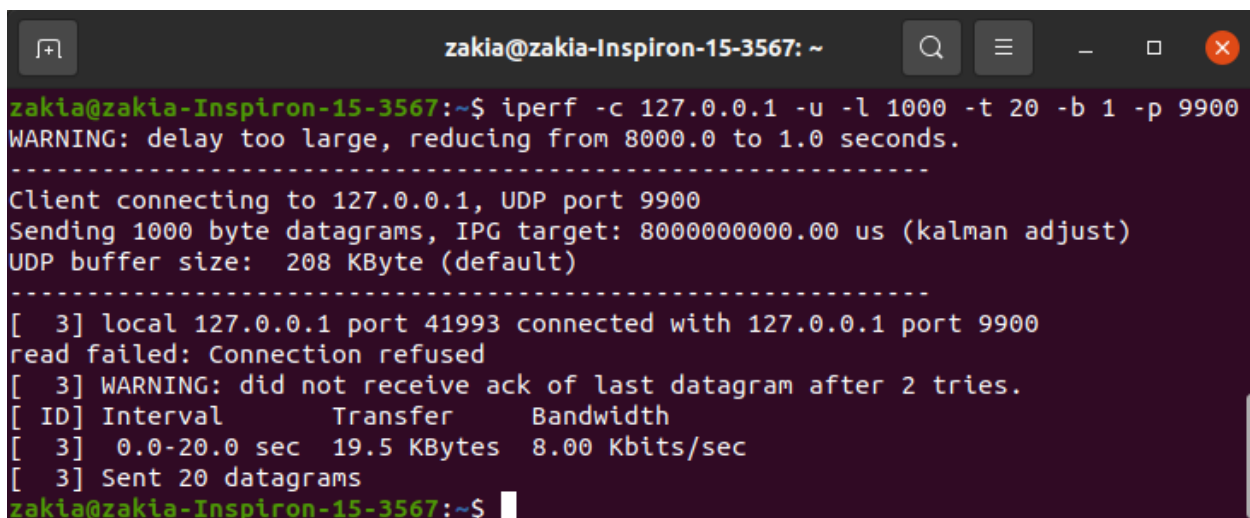
- o Packet length = 1000bytes
- o Time = 20 seconds
- o Bandwidth = 1Mbps
- o Port = 9900

Which are the command lines?

The command lines are:

For terminal 1:

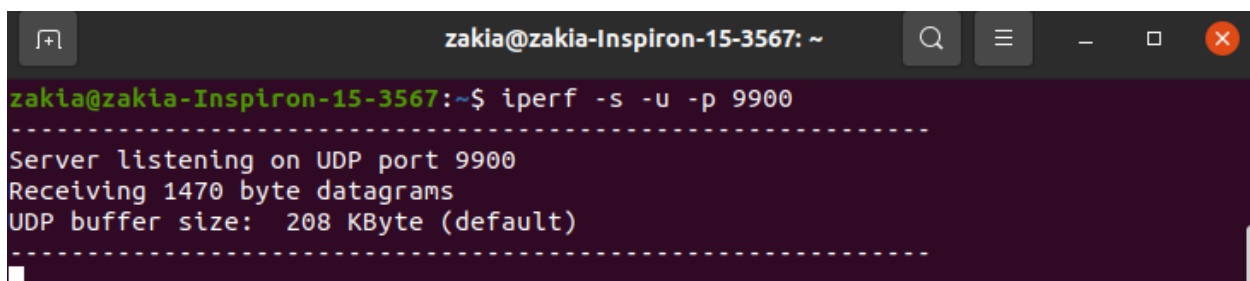
`iperf -c 127.0.0.1 -u -l 1000 -t 20 -b 1 -p 9900`

A terminal window titled 'zakia@zakia-Inspiron-15-3567: ~' showing the execution of the iperf client command. The output includes a warning about a large delay, connection details to 127.0.0.1 port 9900, and a summary of the test results.

```
zakia@zakia-Inspiron-15-3567:~$ iperf -c 127.0.0.1 -u -l 1000 -t 20 -b 1 -p 9900
WARNING: delay too large, reducing from 8000.0 to 1.0 seconds.
-----
Client connecting to 127.0.0.1, UDP port 9900
Sending 1000 byte datagrams, IPG target: 8000000000.00 us (kalman adjust)
UDP buffer size: 208 KByte (default)
-----
[ 3] local 127.0.0.1 port 41993 connected with 127.0.0.1 port 9900
read failed: Connection refused
[ 3] WARNING: did not receive ack of last datagram after 2 tries.
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-20.0 sec   19.5 KBytes 8.00 Kbits/sec
[ 3] Sent 20 datagrams
zakia@zakia-Inspiron-15-3567:~$
```

For terminal 2:

`iperf -s -u -p 9900`

A terminal window titled 'zakia@zakia-Inspiron-15-3567: ~' showing the execution of the iperf server command. The output shows the server listening on UDP port 9900 and receiving datagrams.

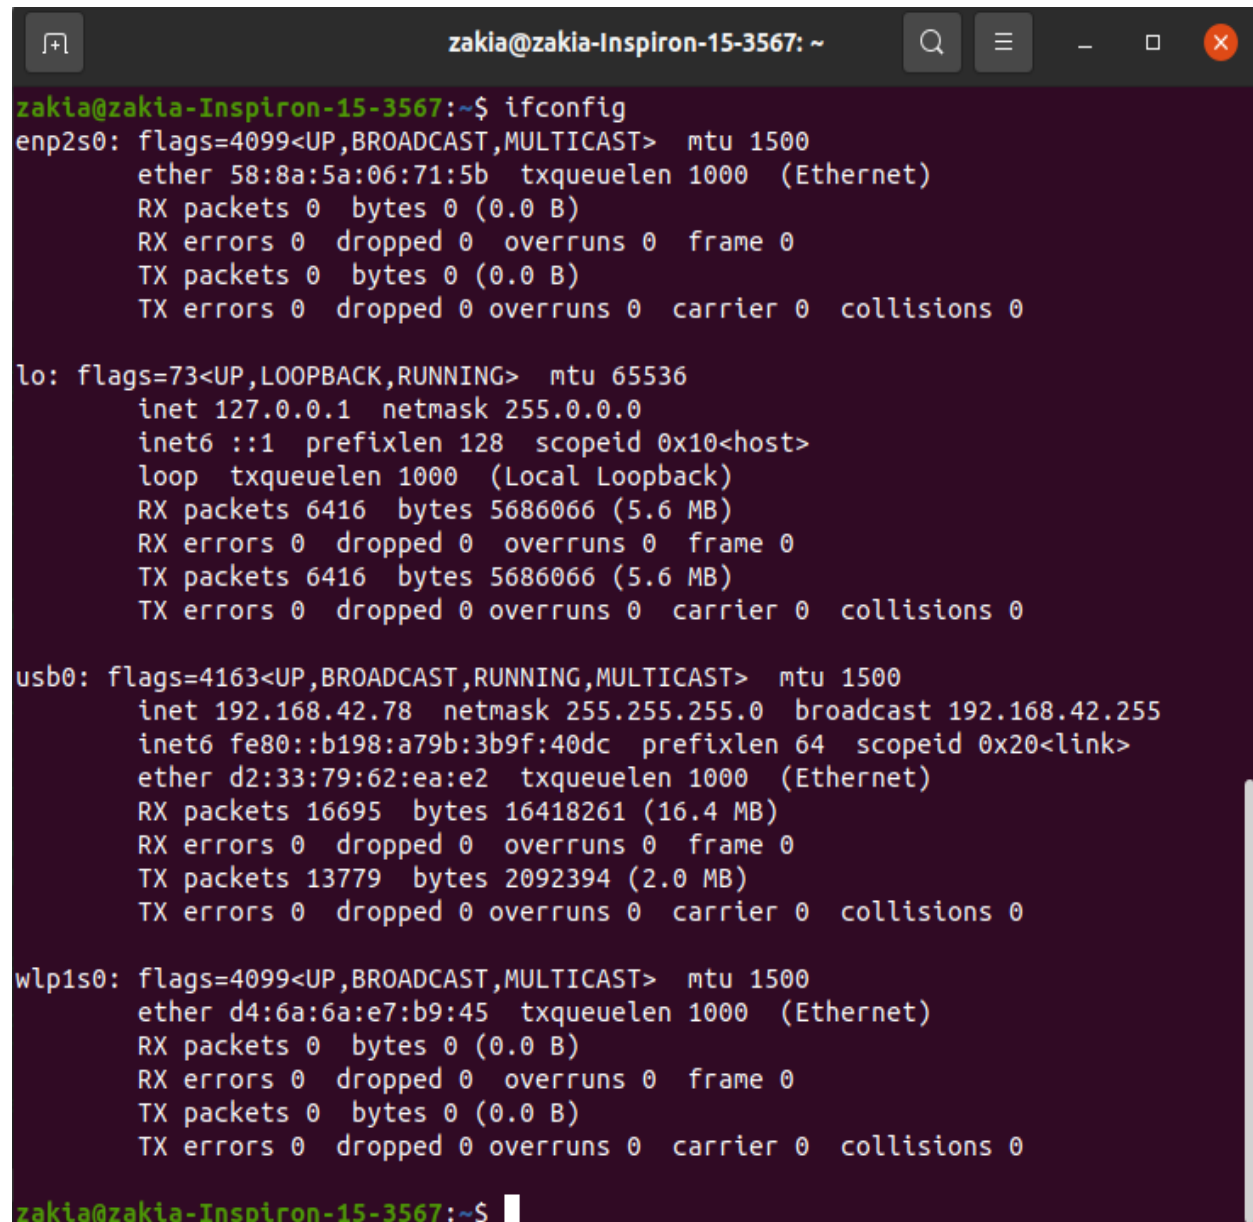
```
zakia@zakia-Inspiron-15-3567:~$ iperf -s -u -p 9900
-----
Server listening on UDP port 9900
Receiving 1470 byte datagrams
UDP buffer size: 208 KByte (default)
-----
```


Using Mininet

Exercise 4.2.1: Open two Linux terminals, and execute the command line `ifconfig` in terminal1. How many interfaces are present?

In terminal-2, execute the command line `sudo mn`, which is the output?

In terminal-1 execute the command line `ifconfig`. How many real and virtual interfaces are present now?



```
zakia@zakia-Inspiron-15-3567: ~  
zakia@zakia-Inspiron-15-3567:~$ ifconfig  
enp2s0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500  
    ether 58:8a:5a:06:71:5b txqueuelen 1000 (Ethernet)  
    RX packets 0 bytes 0 (0.0 B)  
    RX errors 0 dropped 0 overruns 0 frame 0  
    TX packets 0 bytes 0 (0.0 B)  
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0  
  
lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536  
    inet 127.0.0.1 netmask 255.0.0.0  
    inet6 ::1 prefixlen 128 scopeid 0x10<host>  
    loop txqueuelen 1000 (Local Loopback)  
    RX packets 6416 bytes 5686066 (5.6 MB)  
    RX errors 0 dropped 0 overruns 0 frame 0  
    TX packets 6416 bytes 5686066 (5.6 MB)  
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0  
  
usb0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500  
    inet 192.168.42.78 netmask 255.255.255.0 broadcast 192.168.42.255  
    inet6 fe80::b198:a79b:3b9f:40dc prefixlen 64 scopeid 0x20<link>  
    ether d2:33:79:62:ea:e2 txqueuelen 1000 (Ethernet)  
    RX packets 16695 bytes 16418261 (16.4 MB)  
    RX errors 0 dropped 0 overruns 0 frame 0  
    TX packets 13779 bytes 2092394 (2.0 MB)  
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0  
  
wlp1s0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500  
    ether d4:6a:6a:e7:b9:45 txqueuelen 1000 (Ethernet)  
    RX packets 0 bytes 0 (0.0 B)  
    RX errors 0 dropped 0 overruns 0 frame 0  
    TX packets 0 bytes 0 (0.0 B)  
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0  
  
zakia@zakia-Inspiron-15-3567:~$
```



```
zakia@zakia-Inspiron-15-3567:~$ sudo mn
[sudo] password for zakia:
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2
*** Adding switches:
s1
*** Adding links:
(h1, s1) (h2, s1)
*** Configuring hosts
h1 h2
*** Starting controller
c0
*** Starting 1 switches
s1 ...
*** Starting CLI:
mininet>
```

```
zakia@zakia-Inspiron-15-3567: ~  
zakia@zakia-Inspiron-15-3567:~$ ifconfig  
enp2s0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500  
    ether 58:8a:5a:06:71:5b txqueuelen 1000 (Ethernet)  
    RX packets 0 bytes 0 (0.0 B)  
    RX errors 0 dropped 0 overruns 0 frame 0  
    TX packets 0 bytes 0 (0.0 B)  
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0  
  
lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536  
    inet 127.0.0.1 netmask 255.0.0.0  
    inet6 ::1 prefixlen 128 scopeid 0x10<host>  
    loop txqueuelen 1000 (Local Loopback)  
    RX packets 6416 bytes 5686066 (5.6 MB)  
    RX errors 0 dropped 0 overruns 0 frame 0  
    TX packets 6416 bytes 5686066 (5.6 MB)  
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0  
  
usb0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500  
    inet 192.168.42.78 netmask 255.255.255.0 broadcast 192.168.42.255  
    inet6 fe80::b198:a79b:3b9f:40dc prefixlen 64 scopeid 0x20<link>  
    ether d2:33:79:62:ea:e2 txqueuelen 1000 (Ethernet)  
    RX packets 16695 bytes 16418261 (16.4 MB)  
    RX errors 0 dropped 0 overruns 0 frame 0  
    TX packets 13779 bytes 2092394 (2.0 MB)  
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0  
  
wlp1s0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500  
    ether d4:6a:6a:e7:b9:45 txqueuelen 1000 (Ethernet)  
    RX packets 0 bytes 0 (0.0 B)  
    RX errors 0 dropped 0 overruns 0 frame 0  
    TX packets 0 bytes 0 (0.0 B)  
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0  
  
zakia@zakia-Inspiron-15-3567:~$
```

Exercise 4.2.2: Interacting with mininet; in terminal-2, display the following command lines and explain what it does:

mininet> help

```
zakia@zakia-Inspiron-15-3567: ~  
mininet> help  
  
Documented commands (type help <topic>):  
=====
```

EOF	gterm	iperfudp	nodes	pingpair	py	switch
dpctl	help	link	noecho	pingpairfull	quit	time
dump	intfs	links	pingall	ports	sh	x
exit	iperf	net	pingallfull	px	source	xterm

```
  
You may also send a command to a node using:  
  <node> command {args}  
For example:  
  mininet> h1 ifconfig  
  
The interpreter automatically substitutes IP addresses  
for node names when a node is the first arg, so commands  
like  
  mininet> h2 ping h3  
should work.  
  
Some character-oriented interactive commands require  
noecho:  
  mininet> noecho h2 vi foo.py  
However, starting up an xterm/gterm is generally better:  
  mininet> xterm h2  
  
mininet>
```

mininet> nodes

```
zakia...  
mininet> nodes  
available nodes are:  
c0 h1 h2 s1  
mininet>
```

mininet> net

```
zakia@zakia-Inspiron-15-3567: ~  
mininet> net  
h1 h1-eth0:s1-eth1  
h2 h2-eth0:s1-eth2  
s1 lo: s1-eth1:h1-eth0 s1-eth2:h2-eth0  
c0  
mininet> █
```

mininet> dump

```
zakia@zakia-Inspiron-15-3567: ~  
mininet> dump  
<Host h1: h1-eth0:10.0.0.1 pid=4888>  
<Host h2: h2-eth0:10.0.0.2 pid=4890>  
<OVSSwitch s1: lo:127.0.0.1,s1-eth1:None,s1-eth2:None pid=4895>  
<Controller c0: 127.0.0.1:6653 pid=4881>  
mininet> █
```

mininet> h1 ifconfig -a

```
zakia@zakia-Inspiron-15-3567: ~  
mininet> h1 ifconfig -a  
h1-eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500  
    inet 10.0.0.1 netmask 255.0.0.0 broadcast 10.255.255.255  
    inet6 fe80::20c4:ecff:fe5c:574a prefixlen 64 scopeid 0x20<link>  
    ether 22:c4:ec:fc:57:4a txqueuelen 1000 (Ethernet)  
    RX packets 40 bytes 5141 (5.1 KB)  
    RX errors 0 dropped 0 overruns 0 frame 0  
    TX packets 12 bytes 936 (936.0 B)  
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0  
  
lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536  
    inet 127.0.0.1 netmask 255.0.0.0  
    inet6 ::1 prefixlen 128 scopeid 0x10<host>  
    loop txqueuelen 1000 (Local Loopback)  
    RX packets 0 bytes 0 (0.0 B)  
    RX errors 0 dropped 0 overruns 0 frame 0  
    TX packets 0 bytes 0 (0.0 B)  
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0  
  
mininet> █
```

mininet> s1 ifconfig -a

```
zakia@zakia-Inspiron-15-3567: ~  
mininet> s1 ifconfig -a  
enp2s0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500  
    ether 58:8a:5a:06:71:5b txqueuelen 1000 (Ethernet)  
    RX packets 0 bytes 0 (0.0 B)  
    RX errors 0 dropped 0 overruns 0 frame 0  
    TX packets 0 bytes 0 (0.0 B)  
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0  
  
lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536  
    inet 127.0.0.1 netmask 255.0.0.0  
    inet6 ::1 prefixlen 128 scopeid 0x10<host>  
    loop txqueuelen 1000 (Local Loopback)  
    RX packets 6946 bytes 5726953 (5.7 MB)  
    RX errors 0 dropped 0 overruns 0 frame 0  
    TX packets 6946 bytes 5726953 (5.7 MB)  
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0  
  
ovs-system: flags=4098<BROADCAST,MULTICAST> mtu 1500  
    ether 06:83:66:fd:de:4a txqueuelen 1000 (Ethernet)  
    RX packets 0 bytes 0 (0.0 B)  
    RX errors 0 dropped 0 overruns 0 frame 0  
    TX packets 0 bytes 0 (0.0 B)  
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0  
  
s1: flags=4098<BROADCAST,MULTICAST> mtu 1500  
    ether ce:95:30:70:74:41 txqueuelen 1000 (Ethernet)  
    RX packets 0 bytes 0 (0.0 B)  
    RX errors 0 dropped 22 overruns 0 frame 0  
    TX packets 0 bytes 0 (0.0 B)  
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0  
  
s1-eth1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500  
    inet6 fe80::8843:52ff:fe99:304a prefixlen 64 scopeid 0x20<link>  
    ether 8a:43:52:99:30:4a txqueuelen 1000 (Ethernet)  
    RX packets 12 bytes 936 (936.0 B)  
    RX errors 0 dropped 0 overruns 0 frame 0
```

```

s1-eth2: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet6 fe80::942d:3ff:fe52:b6f8 prefixlen 64 scopeid 0x20<link>
    ether 96:2d:03:52:b6:f8 txqueuelen 1000 (Ethernet)
    RX packets 12 bytes 936 (936.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 40 bytes 5141 (5.1 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

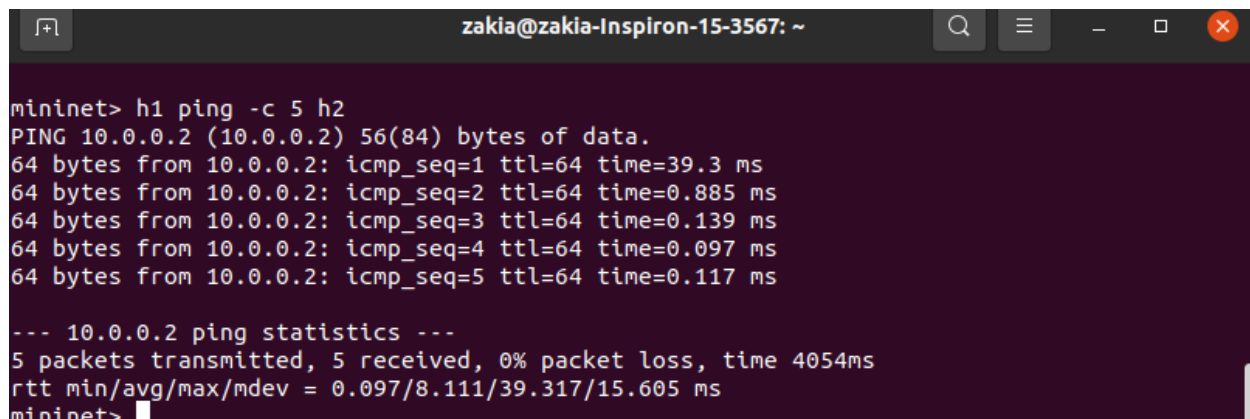
usb0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.42.78 netmask 255.255.255.0 broadcast 192.168.42.255
    inet6 fe80::b198:a79b:3b9f:40dc prefixlen 64 scopeid 0x20<link>
    ether d2:33:79:62:ea:e2 txqueuelen 1000 (Ethernet)
    RX packets 17865 bytes 16613917 (16.6 MB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 15085 bytes 2276072 (2.2 MB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

wlp1s0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    ether d4:6a:6a:e7:b9:45 txqueuelen 1000 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

mininet>

```

mininet> h1 ping -c 5 h2



```

zakia@zakia-Inspiron-15-3567: ~
mininet> h1 ping -c 5 h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data:
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=39.3 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=0.885 ms
64 bytes from 10.0.0.2: icmp_seq=3 ttl=64 time=0.139 ms
64 bytes from 10.0.0.2: icmp_seq=4 ttl=64 time=0.097 ms
64 bytes from 10.0.0.2: icmp_seq=5 ttl=64 time=0.117 ms

--- 10.0.0.2 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4054ms
rtt min/avg/max/mdev = 0.097/8.111/39.317/15.605 ms
mininet>

```

Exercise 4.2.3: In terminal-2, display the following command line: `sudo mn --link tc,bw=10,delay=500ms`

o mininet> h1 ping -c 5 h2, What happen with the link?

o mininet> h1 iperf -s -u &

o mininet> h2 iperf -c IPv4_h1 -u, Is there any packet loss?

```
zakia@zakia-Inspiron-15-3567: ~  
zakia@zakia-Inspiron-15-3567:~$ sudo mn --link tc,bw=10,delay=500ms  
[sudo] password for zakia:  
*** Creating network  
*** Adding controller  
*** Adding hosts:  
h1 h2  
*** Adding switches:  
s1  
*** Adding links:  
(10.00Mbit 500ms delay) (10.00Mbit 500ms delay) (h1, s1) (10.00Mbit 500ms delay) (10.00Mbit 500ms delay) (h2, s1)  
*** Configuring hosts  
h1 h2  
*** Starting controller  
c0  
*** Starting 1 switches  
s1 ... (10.00Mbit 500ms delay) (10.00Mbit 500ms delay)  
*** Starting CLI:  
mininet>
```

```
mininet> h1 ping -c 5 h2  
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.  
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=3035 ms  
64 bytes from 10.0.0.2: icmp_seq=3 ttl=64 time=2017 ms  
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=4068 ms  
64 bytes from 10.0.0.2: icmp_seq=4 ttl=64 time=2000 ms  
64 bytes from 10.0.0.2: icmp_seq=5 ttl=64 time=2000 ms  
  
--- 10.0.0.2 ping statistics ---  
5 packets transmitted, 5 received, 0% packet loss, time 4068ms  
rtt min/avg/max/mdev = 2000.249/2624.156/4068.312/824.792 ms, pipe 4  
mininet>
```

Exercise 4.3.1: Remote controller: The following script reproduce same architecture as before using an external controller: (save as `remote_controller.py`)

Code:


```

1 from mininet.cli import CLI
2 from mininet.net import Mininet
3 from mininet.node import RemoteController
4
5
6 if '__main__' == __name__:
7     net = Mininet(controller=RemoteController)
8     c0 = net.addController('c0', port=6633)
9     s1 = net.addSwitch('s1')
10
11     h1 = net.addHost('h1')
12     h2 = net.addHost('h2')
13
14     net.addLink(s1, h1)
15     net.addLink(s1, h2)
16
17     net.build()
18     c0.start()
19     s1.start([c0])
20     net.startTerms()
21     CLI(net)
22     net.stop()

```

The following scrip defines the behavior of the controller: (save as simple_switch_13.py, it is also available in blackboard)

```
} *simple_switch_13 ☒
```

```
1 from ryu.base import app_manager
2 from ryu.controller import ofp_event
3 from ryu.controller.handler import CONFIG_DISPATCHER, MAIN_DISPATCHER
4 from ryu.controller.handler import set_ev_cls
5 from ryu.ofproto import ofproto_v1_3
6 from ryu.lib.packet import packet
7 from ryu.lib.packet import ethernet
8
9
10 class SimpleSwitch13(app_manager.RyuApp):
11     OFP_VERSIONS = [ofproto_v1_3.OFP_VERSION]
12
13     def __init__(self, *args, **kwargs):
14         super(SimpleSwitch13, self).__init__(*args, **kwargs)
15         self.mac_to_port = {}
16
17     @set_ev_cls(ofp_event.EventOFPSwitchFeatures, CONFIG_DISPATCHER)
18     def switch_features_handler(self, ev):
19         datapath = ev.msg.datapath
20         ofproto = datapath.ofproto
21         parser = datapath.ofproto_parser
22
23         # install table-miss flow entry
24         #
25         # We specify NO BUFFER to max_len of the output action due to
26         # OVS bug. At this moment, if we specify a lesser number, e.g.,
27         # 128, OVS will send Packet-In with invalid buffer_id and
28         # truncated packet data. In that case, we cannot output packets
29         # correctly.
30         match = parser.OFPMatch()
31         actions = [parser.OFPActionOutput(ofproto.OFPP_CONTROLLER,
32                                         ofproto.OFPCML_NO_BUFFER)]
33         self.add_flow(datapath, 0, match, actions)
34     def add_flow(self, datapath, priority, match, actions):
```

```

35     ofproto = datapath.ofproto
36     parser = datapath.ofproto_parser
37
38     inst = [parser.OFPIInstructionActions(ofproto.OFPIT_APPLY_ACTIONS,
39                                           actions)]
40
41     mod = parser.OFPFlowMod(datapath=datapath, priority=priority,
42                             match=match, instructions=inst)
43     datapath.send_msg(mod)
44
45 @set_ev_cls(ofp_event.EventOFPPacketIn, MAIN_DISPATCHER)
46 def _packet_in_handler(self, ev):
47     msg = ev.msg
48     datapath = msg.datapath
49     ofproto = datapath.ofproto
50     parser = datapath.ofproto_parser
51     in_port = msg.match['in_port']
52     pkt = packet.Packet(msg.data)
53     eth = pkt.get_protocols(ethernet.ethernet)[0]
54
55     dst = eth.dst
56     src = eth.src
57
58     dpid = datapath.id
59     self.mac_to_port.setdefault(dpid, {})
60
61     self.logger.info("packet in %s %s %s %s", dpid, src, dst, in_port)
62
63     # learn a mac address to avoid FLOOD next time.
64     self.mac_to_port[dpid][src] = in_port
65
66     if dst in self.mac_to_port[dpid]:
67         out_port = self.mac_to_port[dpid][dst]
68     else:
69         out_port = ofproto.OFPP_FLOOD
70
71     actions = [parser.OFPAActionOutput(out_port)]
72
73     # install a flow to avoid packet_in next time
74     if out_port != ofproto.OFPP_FLOOD:
75         match = parser.OFPMatch(in_port=in_port, eth_dst=dst)
76         self.add_flow(datapath, 1, match, actions)
77
78     data = None
79     if msg.buffer_id == ofproto.OFP_NO_BUFFER:
80         data = msg.data
81
82     out = parser.OFPPacketOut(datapath=datapath, buffer_id=msg.buffer_id,
83                              in_port=in_port, actions=actions, data=data)
84     datapath.send_msg(out)

```

Conclusion:

Mininet is a useful tool for teaching, development and research. With it, a realistic virtual network, running a real kernel switch and application code, can be set up in a few seconds on a single machine, either virtual or native. It is actively developed and supported. Emulation refers to the running of unchanged code on

virtual hardware on the top of the physical host, interactively. It is handy, practical and low cost. It comes with certain restrictions, though, like slower speeds compared to running the same code on a hardware test-bed which is fast and accurate, but expensive. While a simulator requires code modifications and is slow as well. Mininet is a network emulator that enables the creation of a network of virtual hosts, switches, controllers, and links. Mininet hosts standard Linux network software, and its switches support OpenFlow, a software defined network (SDN) for highly flexible custom routing. It constructs a virtual network that appears to be a real physical network. You can create a network topology, simulate it and implement the various network performance parameters such as bandwidth, latency, packet loss, etc, with Mininet, using simple code. You can create the virtual network on a single machine (a VM, the cloud or a native machine). Mininet permits the creation of multiple nodes (hosts, switches or controllers), enabling a big network to be simulated on a single PC. This is very useful in experimenting with various topologies and different controllers, for different network scenarios. The programs that you run can send packets through virtual switches that seem like real Ethernet interfaces, with a given link speed and delay. Packets get processed by what looks like a real Ethernet switch, router, or middle-box, with a given amount of queuing. The Mininet CLI and API facilitate easy interaction with our network. Virtual hosts, switches, links and controllers created through Mininet are the real thing. They are just created using the Mininet emulator rather than hardware and for the most part, their behaviour is similar to discrete hardware elements.

Questions:

Question 5.1: Explain how the traffic generators work?

Answer:

There are a number of ways that traffic generator programs work. Some of them are purely robotic, while others are more detailed SEO-task programs that, while they do what they do very effectively, they don't play by the rules and can earn your site several search penalties. In every case, you should avoid using such programs.

- Robotic refreshing
- Traffic routed through proxy servers
- Spoofed user agents
- Clickfarm traffic
- Hijacked traffic from hacked sites
- Automatic article submission.

Question 5.3: Which is the main difference between configuring UDP and TCP traffic?

Answer:

Transmission control protocol (TCP)	User datagram protocol (UDP)
TCP is a connection-oriented protocol. Connection-orientation means that the communicating devices should establish a connection before transmitting data and should close the connection after transmitting the data.	UDP is the Datagram oriented protocol. This is because there is no overhead for opening a connection, maintaining a connection, and terminating a connection. UDP is efficient for broadcast and multicast type of network transmission.
TCP is reliable as it guarantees delivery of data to the destination router.	The delivery of data to the destination cannot be guaranteed in UDP.
TCP provides extensive error checking mechanisms. It is because it provides flow control and acknowledgment of data.	UDP has only the basic error checking mechanism using checksums.
Sequencing of data is a feature of Transmission Control Protocol (TCP). this means that packets arrive in-order at the receiver.	There is no sequencing of data in UDP. If ordering is required, it has to be managed by the application layer.
TCP is comparatively slower than UDP.	UDP is faster, simpler and more efficient than TCP.
Retransmission of lost packets is possible in TCP, but not in UDP.	There is no retransmission of lost packets in User Datagram Protocol (UDP).
TCP has a (20-80) bytes variable length header.	UDP has a 8 bytes fixed length header.

Question 5.4: In your opinions which are the main advantages and disadvantages of working with mininet? Provide at least 3

Answer:

Advantages:

- **Boots faster:** seconds instead of minutes
- **Scales larger:** hundreds of hosts and switches vs. single digits
- **Provides more bandwidth:** typically 2Gbps total bandwidth on modest hardware
- **Installs easily:** a prepackaged VM is available that runs on VMware or VirtualBox for Mac/Win/Linux with OpenFlow v1.0 tools already installed.

Disadvantages:

Mininet-based networks cannot (currently) exceed the CPU or bandwidth available on a single server.

Mininet cannot (currently) run non-Linux-compatible OpenFlow switches or applications; this has not been a major issue in practice.

Question 5.6: What is the advantage of having a programmable Controller?

Answer:

- many inputs and outputs, excellent for controlling and monitoring many processes.
- designed for industrial environments, very robust and reliable.
- reprogrammable.
- modular.
- ideally suited to supervisory control.
- easy to set up and good for FMS environment.